



Movielens_Scala

default ▾

```
%md
## Working with Spark Dataframes using Scala APIs
```

FINISHED ▶ ⌵ 📖 ⚙️

Working with Spark Dataframes using Scala APIs

Took 0 seconds

```
// read the movielens data file
var ratings = sqlContext.read.format("com.databricks.spark.csv")
    .option("delimiter", "\t").load("file:///home/hadoop/lab/data/movies/ratings.dat")
```

FINISHED ▶ ⌵ 📖 ⚙️

```
ratings: org.apache.spark.sql.DataFrame = [C0: string, C1: string, C2: string, C3: string]
```

Took 76 seconds

```
ratings
```

FINISHED ▶ ⌵ 📖 ⚙️

```
res2: org.apache.spark.sql.DataFrame = [C0: string, C1: string, C2: string, C3: string]
```

Took 1 seconds

```
// display the first 10 records
ratings.show(10)
```

FINISHED ▶ ⌵ 📖 ⚙️

```
+---+---+---+---+
| C0|  C1| C2|      C3|
+---+---+---+---+
|  1|1193|  5|978300760|
|  1| 661|  3|978302109|
|  1| 914|  3|978301968|
|  1|3408|  4|978300275|
|  1|2355|  5|978824291|
|  1|1197|  3|978302268|
```



Notebook ▾

Interpreter

Search in your notebooks



● Connected

MovieLens_Scala

only showing top 10 rows



? ⚙ default ▾

Took 0 seconds

```
// describe the schema of the records
ratings.printSchema()
```

FINISHED ▶ ⌵ 📖 ⚙

```
root
 |-- C0: string (nullable = true)
 |-- C1: string (nullable = true)
 |-- C2: string (nullable = true)
 |-- C3: string (nullable = true)
```

Took 1 seconds

```
ratings.schema
```

FINISHED ▶ ⌵ 📖 ⚙

```
res13: org.apache.spark.sql.types.StructType = StructType(StructField(C0,StringType,true), StructField(C1,StringType,true), StructField(C2,StringType,true), StructField(C3,StringType,true))
```

Took 2 seconds

```
// Applying the schema manually
```

FINISHED ▶ ⌵ 📖 ⚙

```
import org.apache.spark.sql.SQLContext
import org.apache.spark.sql.types.{StructType, StructField, StringType, IntegerType, LongType};
```

```
val ratingSchema = StructType(Array(
  StructField("userid", IntegerType, true),
  StructField("movieid", IntegerType, true),
  StructField("rating", IntegerType, true),
  StructField("timestamp", LongType, true)))
```

```
import org.apache.spark.sql.SQLContext
import org.apache.spark.sql.types.{StructType, StructField, StringType, IntegerType, LongType}
ratingSchema: org.apache.spark.sql.types.StructType = StructType(StructField(userid,IntegerType,true), StructField(movieid,IntegerType,true), StructField(rating,IntegerType,true), StructField(timestamp,LongType,true))
```

Took 2 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
// read the movielens data file with the custom schema
var ratings_df = sqlContext.read.format("com.databricks.spark.csv")
    .option("delimiter", "\t").schema(ratingSchema).load("file:///home/hadoop/lab/data/movies/ratings.dat")
```

ratings_df: org.apache.spark.sql.DataFrame = [userid: int, movieid: int, rating: int, timestamp: bigint]

Took 0 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
ratings_df
```

res21: org.apache.spark.sql.DataFrame = [userid: int, movieid: int, rating: int, timestamp: bigint]

Took 1 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
ratings_df.show( 5 )
```

```
+-----+-----+-----+-----+
|userid|movieid|rating|timestamp|
+-----+-----+-----+-----+
|    1|    1193|    5|978300760|
|    1|     661|    3|978302109|
|    1|     914|    3|978301968|
|    1|    3408|    4|978300275|
|    1|    2355|    5|978824291|
+-----+-----+-----+-----+
only showing top 5 rows
```

Took 1 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
ratings_df.printSchema()
```

```
root
|-- userid: integer (nullable = true)
|-- movieid: integer (nullable = true)
|-- rating: integer (nullable = true)
|-- timestamp: long (nullable = true)
```

Took 0 seconds

```
// Returns a list of columns
ratings_df.columns
```

```
res34: Array[String] = Array(userid, movieid, rating, timestamp)
```

Took 0 seconds

FINISHED ▶ 🔍 📖 ⚙️

```
// How many records in the dataframe?
ratings_df.count()
```

```
res37: Long = 1000209
```

Took 1 seconds

FINISHED ▶ 🔍 📖 ⚙️

```
// We donot need the timestamp column.. let's drop it
ratings_df = ratings_df.drop( "timestamp" )
```

```
ratings_df: org.apache.spark.sql.DataFrame = [userid: int, movieid: int, rating: int]
```

Took 0 seconds

FINISHED ▶ 🔍 📖 ⚙️

```
ratings_df.show(5)
```

```
+-----+-----+-----+
|userid|movieid|rating|
+-----+-----+-----+
|    1|    1193|    5|
|    1|    661|    3|
|    1|    914|    3|
|    1|   3408|    4|
|    1|   2355|    5|
+-----+-----+-----+
only showing top 5 rows
```

Took 1 seconds

FINISHED ▶ 🔍 📖 ⚙️

```
// Lets's group by movieid and find which movie has been rated by how many users
var movie_counts = ratings_df.groupBy("movieid").count()
```

```
movie_counts: org.apache.spark.sql.DataFrame = [movieid: int, count: bigint]
```

Took 0 seconds (outdated)

FINISHED ▶ 🔍 📖 ⚙️

FINISHED ▶ ⌵ 📖 ⚙️

```
import org.apache.spark.sql.functions._
movie_counts = movie_counts.sort(desc("count"))

import org.apache.spark.sql.functions._
movie_counts: org.apache.spark.sql.DataFrame = [movieid: int, count: bigint]
```

Took 1 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
movie_counts.show( 10 )
```

```
+-----+-----+
|movieid|count|
+-----+-----+
|   2858| 3428|
|    260| 2991|
|   1196| 2990|
|   1210| 2883|
|    480| 2672|
|   2028| 2653|
|    589| 2649|
|   2571| 2590|
|   1270| 2583|
|    593| 2578|
+-----+-----+
only showing top 10 rows
```

Took 3 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
// average rating for each movie
var avg_ratings = ratings_df.groupBy("movieid").agg( avg( "rating") )

avg_ratings: org.apache.spark.sql.DataFrame = [movieid: int, avg(rating): double]
```

Took 0 seconds (outdated)

FINISHED ▶ ⌵ 📖 ⚙️

```
avg_ratings.printSchema()
```

```
root
|-- movieid: integer (nullable = true)
|-- avg(rating): double (nullable = true)
```

Took 0 seconds

```
avg_ratings = avg_ratings.sort( desc( "avg(rating)" ) )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
avg_ratings: org.apache.spark.sql.DataFrame = [movieid: int, avg(rating): double]
```

Took 0 seconds

```
avg_ratings.show( 10 )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
+-----+-----+
|movieid|avg(rating)|
+-----+-----+
|   3607|         5.0|
|    989|         5.0|
|    787|         5.0|
|   3172|         5.0|
|   3656|         5.0|
|   3881|         5.0|
|   3382|         5.0|
|   3233|         5.0|
|   3280|         5.0|
|   1830|         5.0|
+-----+-----+
only showing top 10 rows
```

Took 3 seconds

```
var avg_ratings_count = avg_ratings.join( movie_counts, "movieid" )
// An Alternate way of joining two dataframes, in case the column names are different
//var avg_ratings_count = avg_ratings.join( movie_counts, avg_ratings("movieid") === movie_counts("movieid") )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
avg_ratings_count: org.apache.spark.sql.DataFrame = [movieid: int, avg(rating): double, count: bigint]
```

Took 1 seconds

```
avg_ratings_count.show(10)
```

FINISHED ▶ ⌵ 📖 ⚙️

```
+-----+-----+-----+
|movieid|      avg(rating)|count|
+-----+-----+-----+
```

```
+-----+-----+-----+
|      31|3.1134751773049647| 141|
|     231|3.1924242424242424| 660|
|     431|3.6910569105691056| 369|
|     631|                2.08|  75|
|     831|                3.85|  20|
|    1031| 3.479623824451411| 319|
|    1231|                4.08| 750|
|    1431|2.5474137931034484| 232|
|    1631|3.5918367346938775|  49|
|    1831|2.5847076461769114| 667|
+-----+-----+-----+
```

only showing top 10 rows

Took 75 seconds

```
avg_ratings_count.printSchema()
```

FINISHED ▶ ⌵ 📖 ⚙️

```
root
|-- movieid: integer (nullable = true)
|-- avg(rating): double (nullable = true)
|-- count: long (nullable = false)
```

Took 1 seconds

```
// rename the column avg(rating) to mean_rating
avg_ratings_count = avg_ratings_count.withColumnRenamed( "avg(rating)", "mean_rating" )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
avg_ratings_count: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
```

Took 0 seconds (outdated)

```
avg_ratings_count.show( 5 )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
+-----+-----+-----+
|movieid|    mean_rating|count|
+-----+-----+-----+
|      31|3.1134751773049647| 141|
|     231|3.1924242424242424| 660|
|     431|3.6910569105691056| 369|
|     631|                2.08|  75|
```

```
|    831|          3.85|    20|
+-----+-----+-----+
only showing top 5 rows
```

Took 61 seconds

```
// round off the mean_rating value to 2 decimal places
import scala.math.BigDecimal;

// define an udf (sql function) for rounding of
val roundto2places = udf[Double, Double]( BigDecimal(_).setScale(2, BigDecimal.RoundingMode.HALF_UP).toDouble )

// Using the udf on the column to round off all values in the mean_rating column
avg_ratings_count = avg_ratings_count.withColumn( "mean_rating", roundto2places( avg_ratings_count("mean_rating") ) )

import scala.math.BigDecimal
roundto2places: org.apache.spark.sql.UserDefinedFunction = UserDefinedFunction(<function1>,DoubleType,List(DoubleType))
avg_ratings_count: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
```

FINISHED ▶ ⌵ 📖 ⚙️

```
// cache the dataframe to prevent it from re-evaluating it everytime we are going to use it
avg_ratings_count.cache()
avg_ratings_count.show(5)
```

```
res88: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
```

```
+-----+-----+-----+
|movieid|mean_rating|count|
+-----+-----+-----+
|    31|      3.11|  141|
|   231|      3.19|  660|
|   431|      3.69|  369|
|   631|      2.08|   75|
|   831|      3.85|   20|
+-----+-----+-----+
only showing top 5 rows
```

Took 58 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
var avg_ratings_count_sorted = avg_ratings_count.sort( desc( "mean_rating" ) )
```

FINISHED ▶ ⌵ 📖 ⚙️


```
avg_ratings_count_sorted.cache()
avg_ratings_count_sorted.show(5)
```

```
avg_ratings_count_sorted: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
res91: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
```

```
+-----+-----+-----+
|movieid|mean_rating|count|
+-----+-----+-----+
|   3233|         5.0|    2|
|   3656|         5.0|    1|
|   3280|         5.0|    1|
|   3881|         5.0|    1|
|   3172|         5.0|    1|
+-----+-----+-----+
only showing top 5 rows
```

Took 27 seconds

```
// Consider only those movies which are atleast rated by 100 users. The average rating of this movie will be less biased
var avg_ratings_count_more_20 = avg_ratings_count_sorted.filter( avg_ratings_count("count") > 100 )
var popular_movies_ratings = avg_ratings_count_more_20.sort( desc( "mean_rating" ) , desc( "count" ) )
```

FINISHED ▶ ⌘ 📖 ⚙️

```
avg_ratings_count_more_20: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
popular_movies_ratings: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint]
```

Took 1 seconds (outdated)

```
// Read the movies data into a spark dataframe
var movies_df = sqlContext.read.format("com.databricks.spark.csv")
    .option("delimiter", "\t")
    .option("header", "true" )
    .option("inferSchema", "true")
    .load("file:///home/hadoop/lab/data/movies/movies.dat")
```

FINISHED ▶ ⌘ 📖 ⚙️

```
movies_df: org.apache.spark.sql.DataFrame = [movieid: int, name: string, tags: string]
```

Took 2 seconds

```
// cache and display the first 10 records
movies_df.cache()
movies_df.show( 10 )
```

FINISHED ▶ ⌘ 📖 ⚙️

```
res121: org.apache.spark.sql.DataFrame = [movieid: int, name: string, tags: string]
```

```
+-----+-----+-----+
|movieid|      name|      tags|
+-----+-----+-----+
|      1| Toy Story (1995)|Animation|Childre...| |
|      2|   Jumanji (1995)|Adventure|Childre...|
|      3|Grumpier Old Men ...|      Comedy|Romance|
|      4|Waiting to Exhale...|      Comedy|Drama|
|      5|Father of the Bri...|      Comedy|
|      6|      Heat (1995)|Action|Crime|Thri...|
|      7|   Sabrina (1995)|      Comedy|Romance|
|      8| Tom and Huck (1995)|Adventure|Children's|
|      9| Sudden Death (1995)|      Action|
|     10|   GoldenEye (1995)|Action|Adventure|...|
+-----+-----+-----+
```

only showing top 10 rows

Took 1 seconds

```
movies_df.printSchema()
```

FINISHED ▶ ⌵ 📖 ⚙️

```
root
 |-- movieid: integer (nullable = true)
 |-- name: string (nullable = true)
 |-- tags: string (nullable = true)
```

Took 0 seconds

```
var top_movies = popular_movies_ratings.limit(20).join( movies_df, "movieid" )
```

FINISHED ▶ ⌵ 📖 ⚙️

```
top_movies: org.apache.spark.sql.DataFrame = [movieid: int, mean_rating: double, count: bigint, name: string, tags: string]
```

Took 1 seconds

```
top_movies.show()
```

FINISHED ▶ ⌵ 📖 ⚙️

```
+-----+-----+-----+-----+-----+
|movieid|mean_rating|count|      name|      tags|
+-----+-----+-----+-----+-----+
|    2019|      4.56|   628|Seven Samurai (Th...|      Action|Drama|
```

318	4.55	2227	Shawshank Redempt...	Drama
858	4.52	2223	Godfather, The (1...	Action Crime Drama
50	4.52	1783	Usual Suspects, T...	Crime Thriller
745	4.52	657	Close Shave, A (1...	Animation Comedy ...
527	4.51	2304	Schindler's List ...	Drama War
1148	4.51	882	Wrong Trousers, T...	Animation Comedy
922	4.49	470	Sunset Blvd. (a.k...	Film-Noir
1198	4.48	2514	Raiders of the Lo...	Action Adventure
904	4.48	1050	Rear Window (1954)	Mystery Thriller
1178	4.47	230	Paths of Glory (1...	Drama War
260	4.45	2991	Star Wars: Episod...	Action Adventure ...
750	4.45	1367	Dr. Strangelove o...	Sci-Fi War
1212	4.45	480	Third Man, The (1...	Mystery Thriller
1207	4.43	928	To Kill a Mocking...	Drama
720	4.43	438	Wallace & Gromit:...	Animation
3435	4.42	551	Double Indemnity ...	Crime Film-Noir
2762	4.41	2459	Sixth Sense, The ...	Thriller
912	4.41	1669	Casablanca (1942)	Drama Romance War
913	4.4	1043	Maltese Falcon, T...	Film-Noir Mystery

Took 3 seconds

```
// Donot truncate the column values
top_movies.show( truncate = false )
```

FINISHED ▶ ❌ 📖 ⚙️

movieid	mean_rating	count	name	tags
2019	4.56	628	Seven Samurai (The Magnificent Seven) (Shichinin no samurai) (1954)	Action Drama
318	4.55	2227	Shawshank Redemption, The (1994)	Drama
858	4.52	2223	Godfather, The (1972)	Action Crime Drama
50	4.52	1783	Usual Suspects, The (1995)	Crime Thriller
745	4.52	657	Close Shave, A (1995)	Animation Comedy Thriller
527	4.51	2304	Schindler's List (1993)	Drama War
1148	4.51	882	Wrong Trousers, The (1993)	Animation Comedy
922	4.49	470	Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)	Film-Noir
1198	4.48	2514	Raiders of the Lost Ark (1981)	Action Adventure
904	4.48	1050	Rear Window (1954)	Mystery Thriller
1178	4.47	230	Paths of Glory (1957)	Drama War
260	4.45	2991	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Fantasy Sci-Fi

750	4.45	1367	Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb (1963)	Sci-Fi War	
1212	4.45	480	Third Man, The (1949)	Mystery Thriller	
1207	4.43	928	To Kill a Mockingbird (1962)	Drama	
720	4.43	438	Wallace & Gromit: The Best of Aardman Animation (1996)	Animation	
3435	4.42	551	Double Indemnity (1944)	Crime Film-Noir	
2762	4.41	2459	Sixth Sense, The (1999)	Thriller	
912	4.41	1669	Casablanca (1942)	Drama Romance War	
913	4.4	1043	Maltese Falcon, The (1941)	Film-Noir Mystery	
+-----+-----+-----+-----+-----+-----+					

Took 1 seconds (outdated)

FINISHED ▶ ⌵ 📖 ⚙️

```
// Save the top 20 best rated movies in a file
top_movies.write.format("com.databricks.spark.csv")
.option("header", "true")
.save("file:///home/hadoop/lab/results/topmovies")
```

Took 1 seconds

FINISHED ▶ ⌵ 📖 ⚙️

```
%md
```

```
## Exercises
#### Find out 20 worst rated movies. But only consider those movies which are rated by at least 100 users.
#### Find out best 10 and worst 10 movies in the category of action and drams
```

Exercises

Find out 20 worst rated movies. But only consider those movies which are rated by at least 100 users.

Find out best 10 and worst 10 movies in the category of action and drams

Took 2 seconds

READY ▶ ⌵ 📖 ⚙️